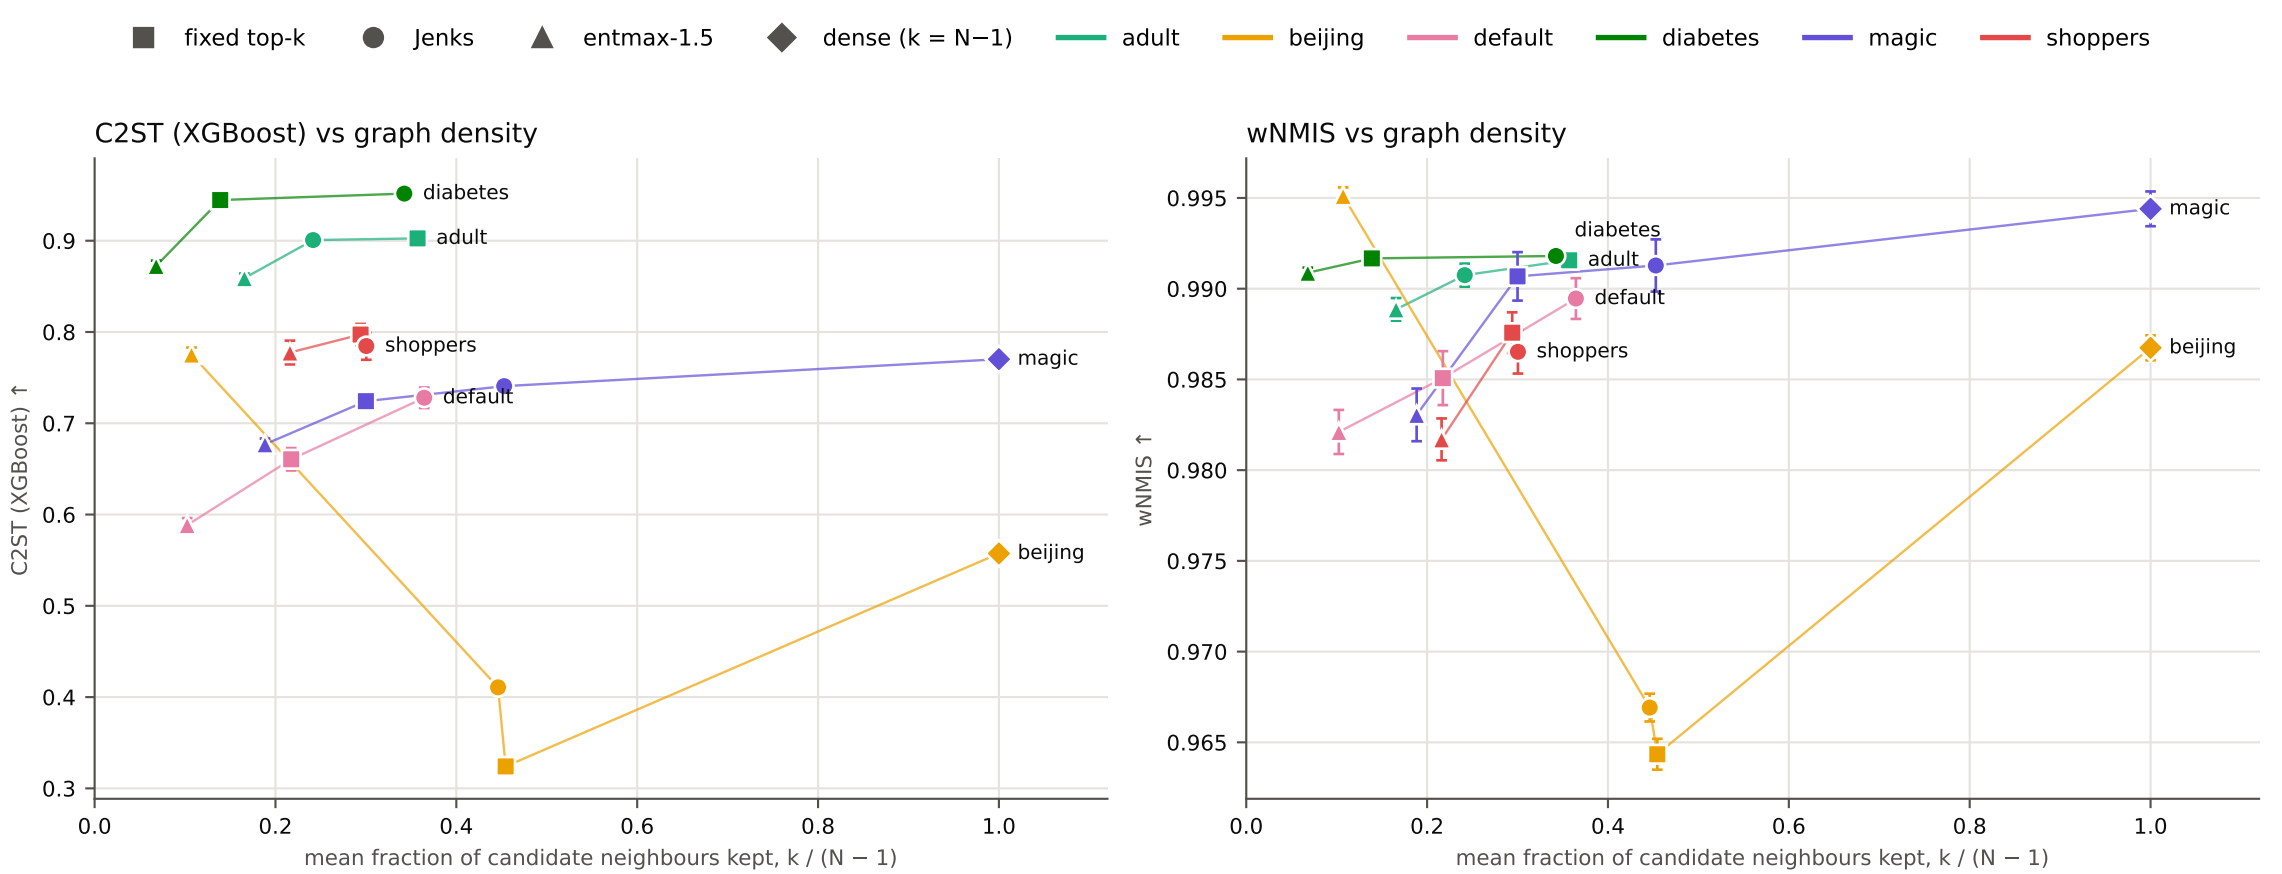

Route A, training seed 0, 20 tables per run (mean ± std, ddof=1); wNMIS and C2ST (XGBoost) from tabpc. x: adaptive rules = mean k over the sampling-phase log (all blocks, t-bins, nodes, count-weighted) / (N − 1); fixed = top_k / (N − 1); dense = 1. Points of one dataset joined in order of x. Not shown (no evaluated run on disk): news.