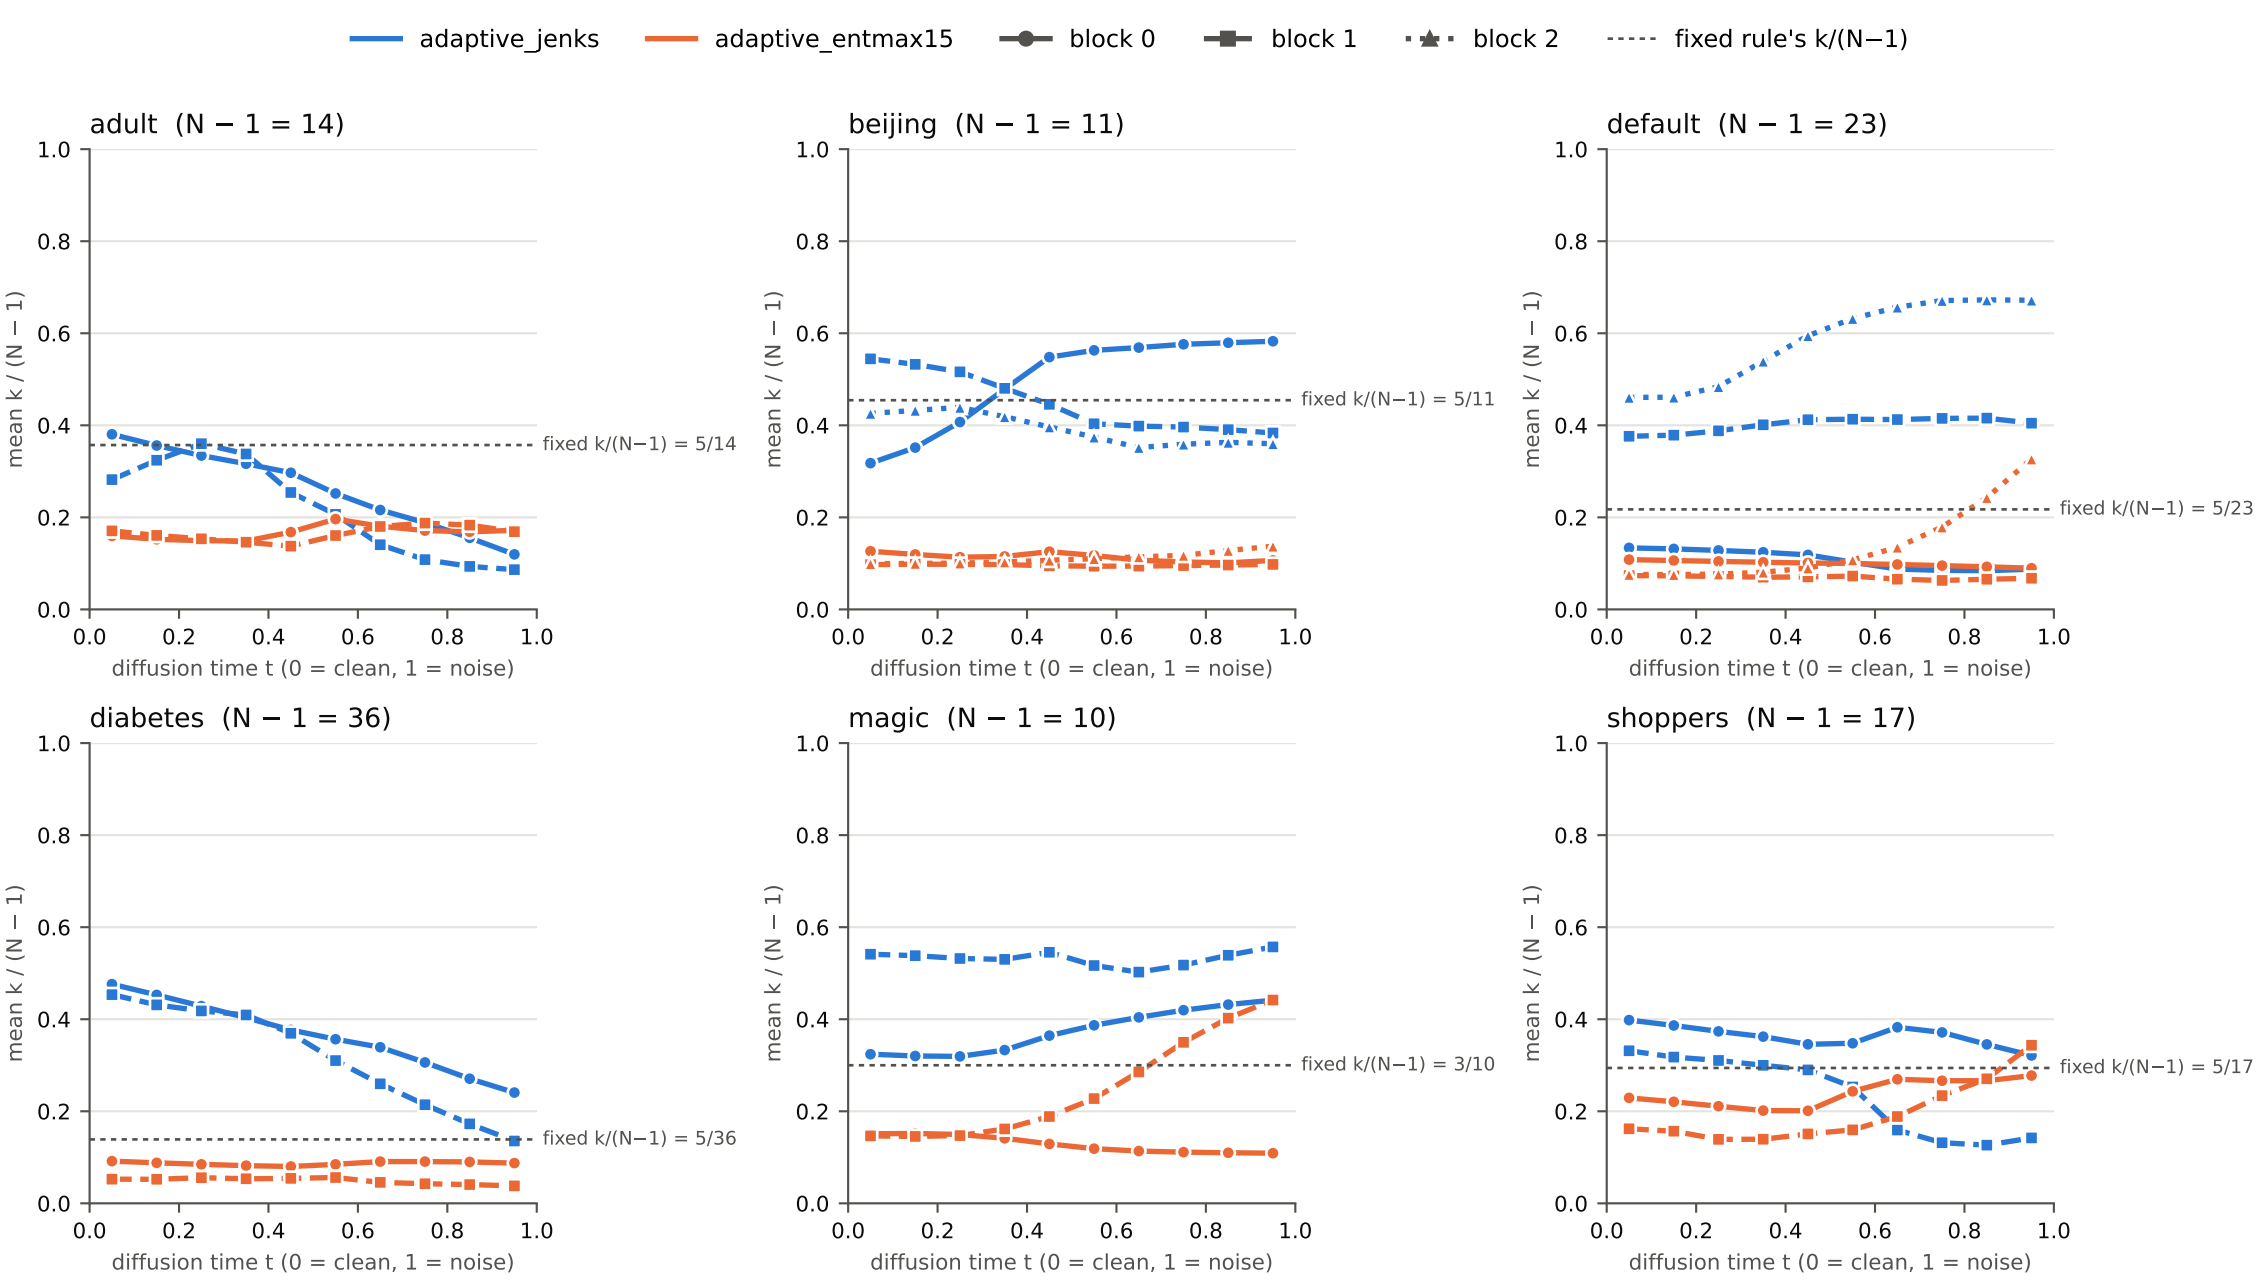

Route A, training seed 0, sampling phase (20 tables); mean over nodes per t-bin of width 0.1. k statistics on disk — jenks: adult, beijing, default, diabetes, magic, shoppers; entmax15: adult, beijing, default, diabetes, magic, shoppers. No adaptive-rule k statistics on disk for news. y = mean k divided by N − 1 (fraction of candidate neighbours kept).