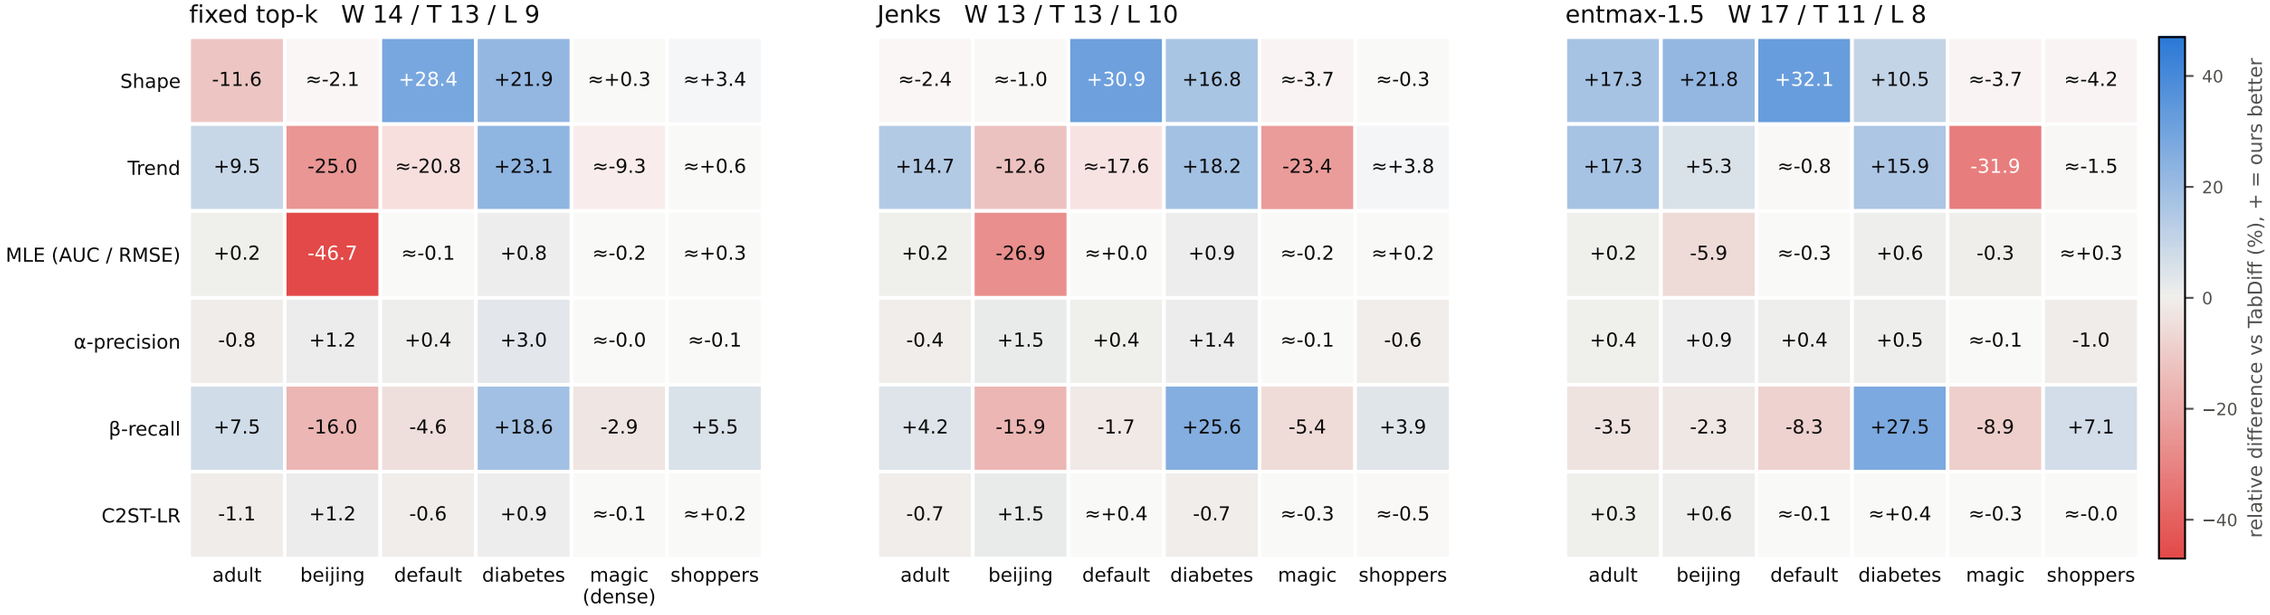

Route A, training seed 0, ours = mean over 20 tables; TabDiff = its paper (Tables 1–3, 7–9). Cell = signed relative difference in %, positive = ours better: lower-is-better metrics (Shape, Trend, MLE on beijing = RMSE) (TabDiff − ours)/TabDiff, higher-is-better (MLE AUC, α, β, C2ST-LR) (ours − TabDiff)/TabDiff. ≈ and light shade = tie, |ours − TabDiff| ≤ our std over the 20 tables (rule of table B6). magic in the fixed panel = the dense run (k = N − 1 = 10), as in the LaTeX tables; table B6 uses the top-k = 3 run (outcomes that change: magic trend: L (B6, topk3) -> T (dense); magic mle: L (B6, topk3) -> T (dense); magic alpha: L (B6, topk3) -> T (dense)). Shared colour scale ±47 % (no clipping). news not shown (no evaluated run on disk).